

HIGH RISK ACTION

AUTHORIZATION RESULT



AUTHORIZED

HUMAN AUTHORIZER



JAMES ANDERSON
Senior Security Analyst



AUTHORIZATION RESULT



DENIED

HUMAN AUTHORIZER



AUTHORIZER UNAVAILABLE
Escalation Required

AUTHORIZATION RESULT



AUTHORIZED

HUMAN AUTHORIZER



ELENA MORRIS
Head of Risk Operations



AUTHORIZATION RESULT



DENIED

HUMAN AUTHORIZER



AUTHORIZER UNAVAILABLE
Escalation Required

EMILIA

Automated control for AI-agent actions.

How EMILIA works · No receipt, no execution.

THE SHIFT

Software stopped asking permission. It just **acts** now.

AI agents already release payments, push deploys, change permissions, curtail datacenters, and command robots — at machine speed, with no human in the loop. The actions are consequential and often **irreversible**. The question every regulator, insurer, and board now asks is the one no one can answer: **who authorized that, and can you prove it?**

THE GAP

**"A human approved this"
is a claim no one can check.**

Today that proof lives in the operator's own **editable logs** — forgeable, backfillable, and un-auditable by a regulator, court, or insurer without trusting the very party under review. Logs tell you what happened. They cannot prove **who authorized it**. That evidence gap is a category — and it is unfilled.

WHAT EMILIA IS

Automated control for every **high-risk agent action**.

Think passport control for AI agents. Before an agent may take a consequential action, EMILIA checks it: is there a valid authorization tracing to a **named human** who approved *this exact action*? If yes — **AUTHORIZED**, proceed. If not — **DENIED**. EMILIA is not the agent's identity and not the tool it calls. It is the checkpoint and the stamp.

AUTHORIZED

named human authorizer present · signed this exact action →
proceed

DENIED

authorizer unavailable / out of scope → **action does not execute**

THE ARTIFACT

A receipt that proves a named human authorized **that exact action**.

TRUST-RECEIPT · VERIFIES OFFLINE

```
action    release_payment $40,000
approver  jchen · Face ID ✓
bound     sha256:9f2c... (exact action)
sig       ed25519:R9w1...
```

✓ verified – no account, no server, no trust in us

No receipt, no execution. Change one byte → verification fails.

HOW IT WORKS

Five steps. **Deny-by-default** until proven.

- 1 **Agent requests a consequential action.** EMILIA intercepts at the actuator boundary and returns **428 — authorization required.**
- 2 **A named human signs the exact action** on their own device (Face ID / passkey). The signature is bound to a hash of the specific action, not a blanket approval.
- 3 **EMILIA verifies** the receipt: valid signature, in-scope, sufficiently assured, not replayed.
- 4 **The action executes** — only now.
- 5 **An execution receipt is emitted** — portable proof it ran, verifiable by anyone, offline.

THE PRODUCT

EMILIA Gate — the **Consequence Firewall**.

Deny-by-default at the actuator boundary. A consequential action runs only with a valid, in-scope, sufficiently-assured, non-replayed receipt — then emits proof it ran. A new security category: not antivirus (is this file bad?), not a firewall (is this packet allowed?) — **is this machine allowed to take this consequence, and can you prove who authorized it?**

Software — MCP tool calls, agent runtimes, CI/CD, systems of record.

Cloud — API-level enforcement in front of any consequential endpoint.

Physical — robots and actuators: no release without a verified receipt.

WHERE IT FITS

Complementary to all. Competitive with none.

The agent-security stack has layers. EMILIA is the one none of the others produce: proof a **named human authorized this exact action**.

IDENTITY / NHI

which machine is acting (the passport)

ACTION AUTH

which tools the agent may call (the visa)

EMILIA

which **named human** authorized **this exact action** – verifiable offline (passport control + the stamp)

Plugs into **MCP**, agent runtimes, **SCITT**, and systems of record.

WHY IT'S TRUSTWORTHY

Verifiable by anyone. Offline. Without trusting us.

A receipt is a self-contained cryptographic object. A regulator, court, insurer, or counterparty can verify it on their own machine — no EMILIA account, no server call, no trust in the operator or in us. Independent verifiers in **JavaScript**, **Python**, and **Go** agree byte-for-byte.

3

independent verifier implementations (JS
· Py · Go)

Offline

no server, no account, no phone-home

One-time

replay-defended — a receipt authorizes
one action

1 byte

change anything → verification fails

OPEN & STANDARDS-BASED

An open protocol, not a **black box**.

EMILIA is Apache-2.0 and built on open standards so it can become shared infrastructure, not a proprietary silo. The reference runtime, verifiers, and conformance vectors are public and reproducible.

Apache-2.0

open reference runtime & verifiers

IETF

authorization-receipts & quorum drafts ·
SCITT signed-statement profile

RR-1

public conformance suite & vectors

4,220

automated tests passing

GET STARTED

Wrap one dangerous action. See it get blocked.

TRY IT OFFLINE · 30 SECONDS

```
npx -y @emilia-protocol/crash-test
```

Spec & quickstart — emiliaprotocol.ai/spec · [/quickstart](https://emiliaprotocol.ai/quickstart)

Live demo — emiliaprotocol.ai/demo · verify a receipt at [/explorer](https://emiliaprotocol.ai/explorer)

Talk to us — team@emiliaprotocol.ai

No receipt, no execution.