EVERY AI ACTION NEEDS AN OWNER

EMILIA PROTOCOL

The trust layer for agentic AI — cryptographic, offline-verifiable proof that a named human approved an irreversible action before it ran.

Seed · 2026 · Confidential — Iman Schrock · team@emiliaprotocol.ai

THE PROBLEM

Agents act. Then no one can say who approved it.

An AI agent releases a payment, changes a vendor's bank details, deletes a record, deploys. Today the only "proof" a human approved it is a **log on the server that did it** — testimony controlled by the party whose conduct is in question.

Not checkable by a third party. Not tamper-evident. As agents move from answering to *acting*, "a human approved this" has to be **evidence**, not testimony.

WHY NOW

The agent-authorization gap just became urgent — and capital is moving.

0.8%

of ~1M agent tool calls are irreversible — and need mandatory human approval (Anthropic autonomy research).

$392M+

into agentic-AI security in the two weeks around RSAC 2026 alone.

all-time high

cyber/AI-security seed valuations in 2025–26, even as the broader market stayed flat.

Every platform is bolting on ad-hoc approval. None of it is portable, interoperable, or independently verifiable. **The standard doesn't exist yet.**

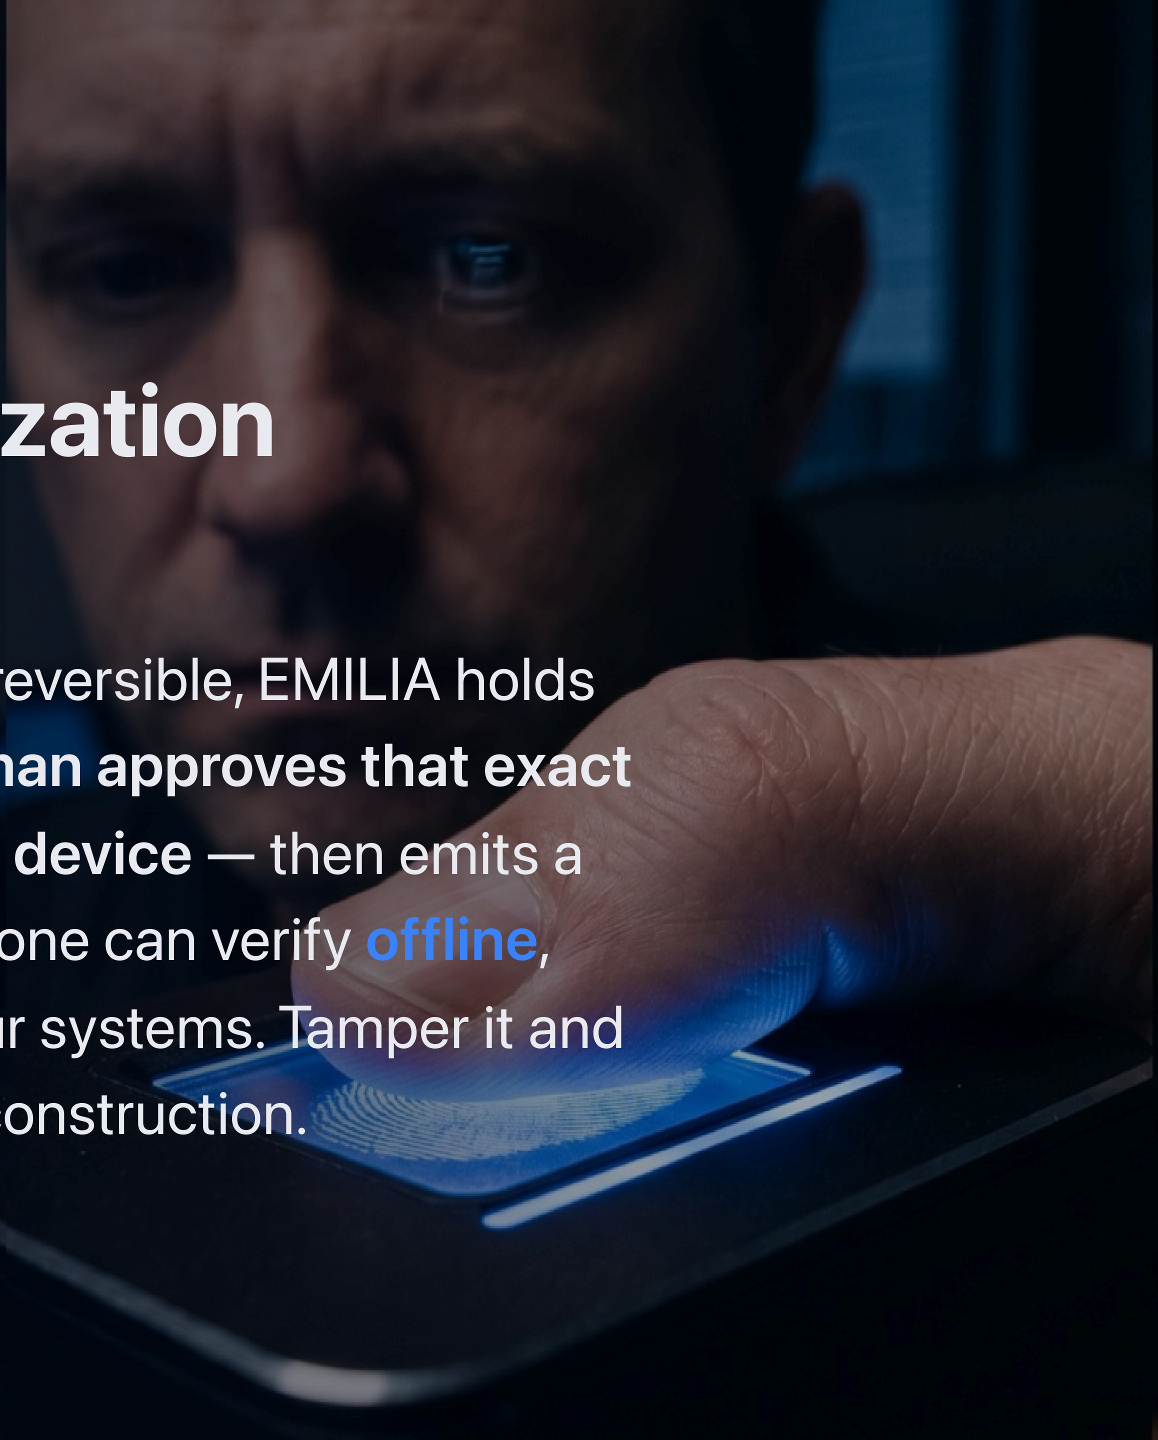

THE SOLUTION

The authorization receipt.

When an action is irreversible, EMILIA holds it until a **named human approves that exact action on their own device** — then emits a portable receipt anyone can verify **offline**, with no access to our systems. Tamper it and verification fails by construction.

HOW IT WORKS

Five stages. One portable proof.

1	2	3	4	5
Eye	**Guard**	**Signoff**	**Receipt**	**Commit**
observes — advisory, never the sole gate	enforces — allow · allow_with_signoff · deny	a named human signs on-device (WebAuthn)	portable, offline-verifiable	Merkle-anchored, sealed

The signed receipt binds the **exact** action (canonical hash) to a **named approver** on a device-bound key. Anyone — auditor, counterparty, regulator — verifies it offline.

WHY IT'S DEFENSIBLE

We're writing the standard, not shipping a plugin.

- → A posted **IETF Internet-Draft** for the receipt format
- → **Three independent verifiers** (JS · Python · Go), npm + PyPI published
- → **Formal verification** (TLA+ & Alloy) enforced in CI
- → **MCP-native** — the agent tool-call boundary, where the ecosystem is converging

Open (Apache-2.0) so it can become the default — monetized by the hosted issuer, registry, and enterprise audit plane on top.

WHY IT COMPOUNDS

One envelope. A registry of profiles. An ecosystem.

Every check is the **same shape** — a signed, offline-verifiable receipt over one envelope — dispatched by a **profile** in an open registry:

payments | provenance (agent hops) | saw = signed | ran = approved | revocation | continuous-eval

A new high-stakes action isn't a new product — it's a new **profile**. **Anyone** can publish one in the reserved namespace, and our verifier enforces it sight-unseen. That's how a toolbox becomes a standard — COSE did it for signing, MIME for content, Sigstore for trust. **The ecosystem is the registry.**

THE WEDGE

Everyone else watches the agent. We make the human provable.

The funded players do **runtime monitoring** of agents. EMILIA does the thing none of them do: **pre-execution, cryptographic, named-human authorization** that survives off their servers.

COMPANY	ROUND	APPROACH
Lyrie.ai	$2M pre-seed	agent trust protocol (exact lane)
Capsule Security	$7M seed (Forgepoint)	runtime "trust layer" for agents
Runlayer	$11M seed (Khosla, Felicis)	MCP security / runtime
EMILIA	raising now	offline-verifiable human authorization, pre-execution

Demand-side, not supply-side.

- → **MCP standard** — a SEP for a consent + receipt hook at the tool-call boundary
- → **x402 / AP2 payment rail** — a shipped HTTP-402 "require-receipt" adapter: an agent-payment client is refused until it presents a valid authorization receipt
- → **Public-sector fraud integrity** — "money moved, no one accountable" is a live, funded pain
- → **The one line:** no receipt, no irreversible action

Win the chokepoint and the receipt becomes table stakes — issuance, registry, and audit are the revenue.

WHAT'S REAL TODAY

Shipped and verifiable — not a demo.

Built & published

IETF I-D posted · `EP-ENVELOPE-v1` + content-addressed profile registry (live) · x402 require-receipt adapter · npm `@emilia-protocol/verify` + Python/Go verifiers · MCP server + middleware · formal proofs in CI · live `/try` Face-ID demo · conformance vectors.

Honest status

Pre-revenue. Zero paying customers. This is a reference implementation + an open standard with a real, formally-verified stack — the raise is to convert that into requiredness and revenue.

Anyone can verify a real receipt offline in ~60 seconds: `npx @emilia-protocol/issue demo` · `emiliaprotocol.ai/try`

TEAM

Founder-led, standards-deep — and hiring the gap.

Iman Schrock — founder. Designed and shipped the protocol, the verifiers, the formal models, and the IETF draft solo.

I'm honest that a solo founder is the main thing to de-risk. **First use of this round** is a technical co-founder / founding engineer and one design partner — turning a one-person standard into a team-backed company.

THE ASK

Raising $2M on a SAFE · ~$20M post-money cap.

Use of funds

1) Technical co-founder + founding engineer.
2) 1–2 design partners (gov fraud-integrity, agent platforms).
3) Hosted issuer / registry / audit plane — the commercial wrapper.
4) Drive the MCP + IETF standardization.

Milestones this round buys

→ Co-founder onboarded.
→ First design partner live on the verifiers.
→ Consent+receipt hook accepted into an ecosystem standard.
→ First paid hosted-issuer pilot.

Priced on the standards moat and the category — not the seed median. EMILIA is **more built than most seed-stage companies**: a posted IETF standard, three independent verifiers, and formal verification, already shipped. Comps: Capsule ($7M seed), Runlayer ($11M seed), Helmet ($9M) — EMILIA goes deeper on the protocol itself.

EMILIA PROTOCOL

No irreversible agent action without a verifiable receipt.

Iman Schrock · team@emiliaprotocol.ai · emiliaprotocol.ai