EMILIA PROTOCOL

The seatbelt for the agentic era.

Verifiable human authorization for the irreversible actions AI agents are starting to take — an open standard, verifiable offline by anyone.

Seed — 2026

team@emiliaprotocol.ai

THE PROBLEM

Fifty years of security asked "who is allowed in?"

The dominant users of software are no longer humans — they're AI agents that move money, change records, and delete data on the fly. The unanswered question is no longer authentication. It's authorization at the moment of action:

Who approved that exact action — before it executed?

AI now creates actions faster than humans can audit them.

— An agent approves a vendor · swaps a bank account · routes benefits · rewrites a record · triggers procurement — in milliseconds, unattended, at scale.

— Every evidence system we have — logs, approvals, audit trails — was built for the speed of *humans* taking those actions.

— Machine-speed action, human-speed audit: that mismatch is the discontinuity. "A human was in the loop" is now an unfalsifiable claim with no artifact behind it.

WHAT EP IS

Authorization receipts.

Before a high-risk action runs, a **named human** approves the *exact* action on their own device (WebAuthn / Face ID). EP emits a signed evidence packet anyone can verify offline, with no account, without trusting our servers.

Alter one byte and verification fails. The receipt stays valid even if EMILIA disappears. `npx @emilia-protocol/crash-test`

SIGNALS OF EXTERNAL VALIDATION

People outside the project have started to engage.

California Dept. of Technology

EP is in the state's GenAI vendor-evaluation process for human-oversight controls.

CA legislative committee

Staff engaged on verifiable-oversight requirements; follow-up set.

First independent implementer

An outside engineer re-ran the verifier + conformance and confirmed it — the first external reliance event.

On the record

Authorization-receipt + quorum IETF drafts posted; MCP server live on the official registry, Glama & Smithery.

Under the hood it is real — 3 independent verifiers (JS/Python/Go) over 8 conformance suites, machine-checked TLA+/Alloy. But the headline is that the market has started to respond.

HOW IT BECOMES INEVITABLE

The reliance ladder.

Legislative engagement ✓

CA committee staff + Dept. of Technology evaluation — engaged now.

Technical validation ✓

First independent implementer verified the receipts — done (2026).

Design partner → now

First org runs gated actions on the verifiers. This round funds it.

Auditor reliance → 2026

An auditor relies on an EP receipt as evidence — the venture-grade milestone.

Insurer acceptance → 2027

Coverage conditioned on receipts for wire / benefit fraud.

Ecosystem standard → 2027+

Receipt expected by default; its absence is the anomaly.

Each rung makes the next cheaper. We're two rungs in — the round buys the third.

FIRST WEDGE — NOT THE WHOLE COMPANY

We enter where the rule already exists.

EP is horizontal (gov, finance, healthcare). We **lead with healthcare** because the two-person rule is *already mandated practice* there — the ISMP / Joint Commission independent double-check. So we change no behavior; we just digitize the evidence. PHI-free by construction.

Entry via healthcare-AI vendors + procurement platforms (B2B2H). The same receipt then carries into finance and government — the wedge, not the ceiling.

Where high-risk actions live.

Government

Fraud, benefit routing, operator overrides. A $500M state fraud finding engaged.

Finance

Payment-change / beneficiary fraud, treasury disbursements, privileged approvals.

Healthcare

High-alert med double-check, capital procurement, Article 14 medical AI.

Agents

Execution controls for autonomous and agentic systems across all of the above.

BUSINESS MODEL — OPEN CORE

The standard is free. The trust root is the business.

— **Free & permissionless:** the format and verifier — that creates ubiquity.

— **Revenue layer:** the operated trust root — enrollment / approver directory, federated transparency log, revocation, compliance-grade evidence pipelines — that an organization cannot casually self-host.

— **Plus:** hosted verification, sector policy packs, audit & evidence tooling, enterprise support.

WHY INEVITABLE

Precedent → recommendation → requirement.

Public-key crypto existed two decades before the browser padlock made its *absence* unacceptable. EP becomes obvious the same way: when a party with leverage makes the **absence of a receipt** unacceptable —

— an insurer's coverage condition on wire / benefit fraud,

— an auditor's controls framework,

— a government AI-auditor registry that needs something verifiable to audit against.

Insurance turns the thesis into a paying pilot.

Cyber and crime policies already require dual authorization and out-of-band verification as **conditions precedent** — and deny claims when they can't be proven. Deepfakes just broke the callback; California AB 316 pins liability on the deployer.

So a carrier — or an insured's auditor — accepting an EP receipt is two things at once: a paid pilot, and the **auditor-reliance milestone (rung 4)**. EP-QUORUM is a cryptographic two-person rule a cloned voice can't defeat.

In motion: an insurer one-pager, a dedicated page, and outreach to At-Bay, Klaimee, Newfront, Beazley, Corvus.

Raising $2M on a SAFE · ~$20M post-money cap.

Use of funds

1) Technical co-founder + founding engineer. 2) 1–2 design partners (gov fraud-integrity, agent platforms). 3) Hosted issuer / registry / audit plane — the commercial wrapper. 4) Drive MCP + IETF standardization.

Milestones this round buys

Co-founder onboarded · first design partner live on the verifiers · consent+receipt hook accepted into an ecosystem standard · first paid hosted-issuer pilot.

Priced on the standards moat and the category — not the seed median. A posted IETF standard, three independent verifiers, and formal verification, already shipped. Comps: Capsule ($7M seed), Runlayer ($11M seed), Helmet ($9M).

EMILIA PROTOCOL

No irreversible agent action without a verifiable receipt.

team@emiliaprotocol.ai · emiliaprotocol.ai · open source (Apache-2.0)