

EMILIA PROTOCOL

The seatbelt for the agentic era.

Verifiable human authorization for the irreversible actions AI agents are starting to take — an open standard, verifiable offline by anyone.

THE PROBLEM

Fifty years of security asked "who is allowed in?"

The dominant users of software are no longer humans — they're AI agents that move money, change records, and delete data on the fly. The unanswered question is no longer authentication. It's **authorization at the moment of action:**

Who approved **that exact action** — before it executed?

AI now creates actions **faster than humans can audit them.**

- An agent approves a vendor · swaps a bank account · routes benefits · rewrites a record · triggers procurement — in milliseconds, unattended, at scale.
- Every evidence system we have — logs, approvals, audit trails — was built for the speed of *humans* taking those actions.
- Machine-speed action, human-speed audit: that mismatch is the discontinuity. "A human was in the loop" is now an **unfalsifiable claim** with no artifact behind it.

WHAT EP IS

Authorization receipts.

Before a high-risk action runs, a **named human** approves the *exact* action on their own device (WebAuthn / Face ID). EP emits a signed evidence packet anyone can verify **offline, with no account, without trusting our servers.**

Alter one byte and verification fails. The receipt stays valid even if EMILIA disappears. **`npx @emilia-protocol/crash-test`**

People outside the project have started to engage.

California Dept. of Technology

EP is in the state's GenAI vendor-evaluation process for human-oversight controls.

First independent implementer

An outside engineer re-ran the verifier + conformance and confirmed it — the first external reliance event.

Under the hood it is real — 3 independent verifiers (JS/Python/Go) over 8 conformance suites, machine-checked TLA+/Alloy. But the headline is that the market has started to respond.

CA legislative committee

Staff engaged on verifiable-oversight requirements; follow-up set.

On the record

Authorization-receipt + quorum IETF drafts posted; MCP server live on the official registry, Glama & Smithery.

HOW IT BECOMES INEVITABLE

The reliance ladder.

Legislative engagement ✓

CA committee staff + Dept. of Technology evaluation — engaged now.

Design partner → now

First org runs gated actions on the verifiers. This round funds it.

Insurer acceptance → 2027

Coverage conditioned on receipts for wire / benefit fraud.

Each rung makes the next cheaper. We're two rungs in — the round buys the third.

Technical validation ✓

First independent implementer verified the receipts — done (2026).

Auditor reliance → 2026

An auditor relies on an EP receipt as evidence — the venture-grade milestone.

Ecosystem standard → 2027+

Receipt expected by default; its absence is the anomaly.

We enter where the rule already exists.

EP is horizontal (gov, finance, healthcare). We lead with **healthcare** because the two-person rule is *already mandated practice* there — the ISMP / Joint Commission independent double-check. So we change no behavior; we just **digitize the evidence**. PHI-free by construction.

Entry via healthcare-AI vendors + procurement platforms (B2B2H). The same receipt then carries into finance and government — the wedge, not the ceiling.

Where high-risk actions live.

Government

Fraud, benefit routing, operator overrides. A \$500M state fraud finding engaged.

Healthcare

High-alert med double-check, capital procurement, Article 14 medical AI.

Finance

Payment-change / beneficiary fraud, treasury disbursements, privileged approvals.

Agents

Execution controls for autonomous and agentic systems across all of the above.

The standard is free. The trust root is the business.

- **Free & permissionless:** the format and verifier — that creates ubiquity.
- **Revenue layer:** the operated trust root — enrollment / approver directory, federated transparency log, revocation, compliance-grade evidence pipelines — that an organization cannot casually self-host.
- **Plus:** hosted verification, sector policy packs, audit & evidence tooling, enterprise support.

Precedent → recommendation → requirement.

Public-key crypto existed two decades before the browser padlock made its *absence* unacceptable. EP becomes obvious the same way: when a party with leverage makes the **absence of a receipt** unacceptable —

- an insurer's coverage condition on wire / benefit fraud,
- an auditor's controls framework,
- a government AI-auditor registry that needs something verifiable to audit against.

Raising \$2M on a SAFE · ~\$20M post-money cap.

Use of funds

1) Technical co-founder + founding engineer. 2) 1–2 design partners (gov fraud-integrity, agent platforms). 3) Hosted issuer / registry / audit plane — the commercial wrapper. 4) Drive MCP + IETF standardization.

Priced on the standards moat and the category — not the seed median.

A posted IETF standard, three independent verifiers, and formal verification, already shipped. Comps: Capsule (\$7M seed), Runlayer (\$11M seed), Helmet (\$9M).

Milestones this round buys

Co-founder onboarded · first design partner live on the verifiers · consent+receipt hook accepted into an ecosystem standard · first paid hosted-issuer pilot.

EMILIA PROTOCOL

No irreversible agent action without a verifiable receipt.

team@emiliaprotocol.ai · emiliaprotocol.ai · open source

(Apache-2.0)