The recourse layer for autonomous actions

Agents now move money, change payees, deploy code, delete records, and sign commitments — a new loss vector on your book. You can't price it, because the one fact a claim turns on — who authorized this exact action — is today self-reported, stored in the deploying party's own mutable log, and unverifiable by the party paying the claim.

> **EMILIA is the artifact that makes an agent-action loss underwritable: proof, verifiable by you, offline, of who authorized this exact action before it ran.**

THE GAP YOU'RE HITTING

- Fraud lives **inside a valid session** (business-email-compromise-style payee swaps): identity and access controls all pass; nothing proves a human approved the exact payment.
- The "human approved it" record is a click in the insured's own database — the party under examination. Not evidence you can rely on.
- So agent-action risk is either **denied, excluded, or priced blind.** Travelers v. ICS-style disputes turn on exactly this missing artifact.

WHAT EMILIA GIVES THE CARRIER

- A per-action, **offline-verifiable** authorization receipt: a named human (or an m-of-n quorum) device-bound-signed the **exact canonical action**, before execution, one-time, tamper-evident.
- **Execution binding** — the amount/beneficiary/target that ran must match what was authorized (catches the $250K-to-A → $300K-to-B swap).
- **Verdict-complete denials** — a refused or absent approval is itself a signed, provable event. "The human was never asked" becomes evidence, not a gap.
- Retention + a status/revocation rail for the evidence a claim needs.

THE INSTRUMENT: RECOURSE REFERENCE

You issue a signed **Recourse Reference** — responsible entity, coverage class + limit, exclusions (bound by hash), dispute endpoint, evidence requirements, status URL — **bound by digest to the authorization receipt.** A gateway checks it before the action runs:

```
Authorization: yes
Identity: yes
Receipt: yes (named human, exact action)
Recourse: <your signed reference + terms + status>
```

A claim presents the receipt + execution attestation as the evidence **you specified**. EMILIA proves the commitment and the authorization; **you bear the loss and adjudicate the exclusions.** Runnable today: a deterministic Recourse Reference vector with all five fail-closed refusal cases (tampered terms, wrong action, unbound authorization, expired window, unpinned issuer).

THE ASK — A BOUNDED PILOT LINE

Co-underwrite one action class against EMILIA receipts

- Pick one line: vendor-payment release, or payee/bank-detail change.
- We co-define the evidence requirements + the exclusions binding.
- You price a small, capped liability limit against verified receipts.
- EMILIA supplies reference issuance, verification, retention, and the runnable vectors — no new crypto, Apache-2.0.

Honest boundary. EMILIA proves the authorization and the recourse commitment, and binds them to the exact action — verifiable by you, offline, independent of the insured's systems. EMILIA does **not** bear the loss, decide whether a loss falls inside your exclusions, verify the responsible party's solvency, or move funds — those are yours. EMILIA's contribution is that none of that adjudication can even begin without a verifiable, action-bound, authorization-backed artifact, and that is exactly what it produces. Necessary, not sufficient — and today, missing.

Maturity: runnable Recourse Reference + authorization-receipt vectors (Ed25519 over RFC-8785 JCS, zero new cryptography) · three cross-language verifiers agreeing byte-for-byte · IETF Internet-Drafts (authorization receipts; multi-party quorum) · formally modeled (TLA+, Alloy) · open standard, Apache-2.0 — receipts verify even if the vendor disappears.

team@emiliaprotocol.ai
emiliaprotocol.ai · EMILIA Protocol, Inc.