

Making "human oversight" verifiable, not just required.

A drafting note for AI-accountability legislation. The hard part isn't requiring a human in the loop — it's writing the requirement so that, after a consequential automated decision, the state can **prove** a named human actually authorized it. This note describes the property to require, vendor-neutral model language, and how it maps to NIST and the EU AI Act.

THE DRAFTING GAP

Most "meaningful human review" or "human oversight" provisions are **unfalsifiable as written**: an agency or vendor can assert a human reviewed a decision, and there is no artifact a court, auditor, or affected person can independently check. The fix is to require that the oversight **produce verifiable evidence** — not to specify a product or technology.

FROM AN UNVERIFIABLE CLAUSE TO A VERIFIABLE ONE

COMMON DRAFTING (UNFALSIFIABLE)	VERIFIABLE-EVIDENCE DRAFTING
"A human shall review consequential automated decisions before they take effect."	...and shall generate a record identifying the named individual who approved the specific action , verifiable independently of the system that produced it.
"The agency shall maintain logs of human review."	Logs are controlled by the reviewed party. Require evidence that is tamper-evident and checkable offline by a third party without trusting that party's systems.
"Oversight consistent with [named vendor / product]."	Require a property (verifiable authorization evidence), not a vendor — preserving competition and avoiding capture or obsolescence.

MODEL LANGUAGE YOU CAN ADAPT (TECHNOLOGY-NEUTRAL)

ILLUSTRATIVE — FOR STAFF TO ADAPT, NOT STATUTORY TEXT

"Before a [high-risk automated decision] takes effect, a covered entity shall obtain authorization from a **named natural person** and shall produce a **verifiable authorization record** that (a) identifies the approver, (b) is cryptographically bound to the specific action and its parameters, (c) records the time of approval relative to execution, and (d) can be **verified by an independent party using openly available means, without access to or trust in the covered entity's systems**. The record shall be retained and made available to [the auditor / the affected individual on request]."

WHY "OPEN AND INDEPENDENTLY VERIFIABLE" BELONGS IN THE TEXT

- **Avoids vendor lock-in and capture** — any conforming implementation satisfies it; the state never depends on one company.
- **Survives an appeal** — an affected person (or their counsel) can verify the authorization themselves, not take the agency's word.
- **Future-proof** — it specifies an outcome (checkable evidence), so it doesn't go stale as tools change.
- **Already implementable** — open standards that produce exactly this evidence exist today; the requirement is not aspirational.

ALIGNMENT WITH EXISTING FRAMEWORKS

This is the operational form of the human-oversight expectations already in the **NIST AI Risk Management Framework** (GOVERN/MAP human-accountability functions) and **EU AI Act Article 14** (human oversight of high-risk AI). Those name the obligation; verifiable authorization evidence is what makes it *auditable* rather than declarative.

What this evidence proves — and doesn't. A verifiable authorization record proves, independently, that a named person approved a specific action under a stated policy and that the record is unaltered. It does not by itself prove the decision was *correct*, nor replace the agency's underlying systems. Whether approvals are genuinely considered (not rubber-stamped) is a process question the evidence makes *auditable*, not one it resolves alone. As a working reference, the open-source **EMILIA Protocol** (Apache-2.0; IETF Internet-Draft *draft-schrock-ep-authorization-receipts*; verifiers in three languages) implements this property today — offered only to show the requirement is concrete and achievable, not as a product to specify.